

Bredariol Francesco



UNIVERSITÀ
DEGLI STUDI
DI TRIESTE

EVES

Evolutionary Elo System

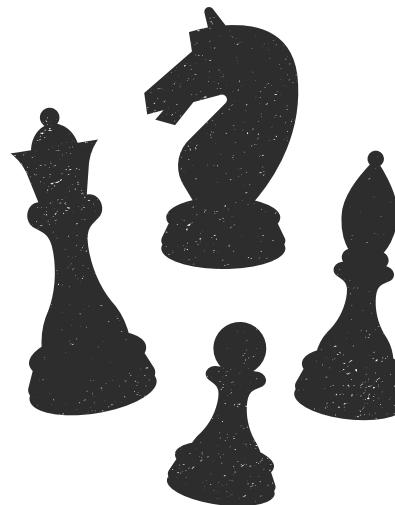
An optimization techniques for zero sum games

The problem

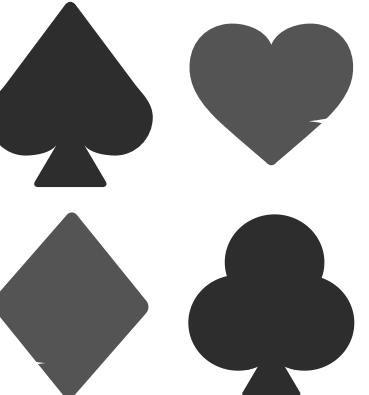
Zero Sum Games

The game theory describes zero-sum game as situations where gains and losses add up to zero.

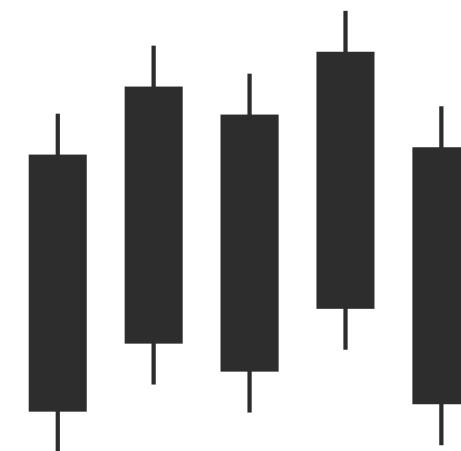
There could be two or more players involved.



chess



poker



futures trading

The problem

Zero Sum Games

While optimizing zero sum games there are some non-trivial tasks to face
(other than the optimization task itself):

- compute how an individual perform
- select fair matches for the individual
- obtain a variety of solutions among the entire population
- map solution into an “human-like” level

The idea

Elo Rating System

The Elo rating system is a method for calculating the relative skill levels of players in zero-sum games.

It is a statistical estimation of the probability of winning based on the historical performances of an individual among the others

It can be used to match individuals with similar skills in order to obtain the most fair possible matches.



Arpad Elo
inventor of the Elo rating system

Elo Rating System

Winning probability

Given an initial population of n players, each of whom has an **Elo level** γ , the probability for the i -th individual of winning against the j -th individual is computed using a **sigmoid function***:

$$p_i(\text{win}|x_j) = \frac{1}{1 + e^{\frac{\gamma_i - \gamma_j}{\lambda}}}$$

Elo Ranking System

Update rule

In a zero-sum games an individual can either win or lose (we assume no tie).

Encoding the result r of a match as 1 for win and 0 for lose, the elo update rule* is the following (from the perspective of the i -th individual against the j -th one):

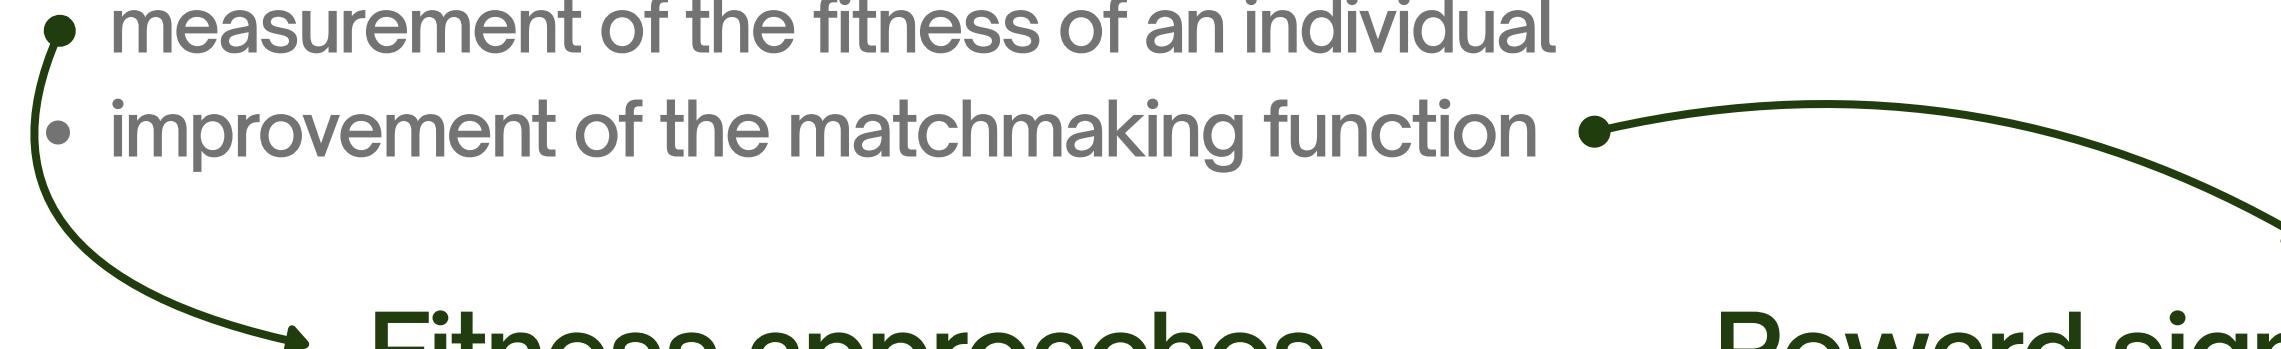
$$\gamma_i = \gamma_i + k(r - p_i(r = 1|x_j))$$

On the way to the EvES

Online/Offline optimization

The Elo described until now is good for two main reasons:

- measurement of the fitness of an individual
- improvement of the matchmaking function



Fitness approaches

evolutionary strategies
genetic algorithms

...

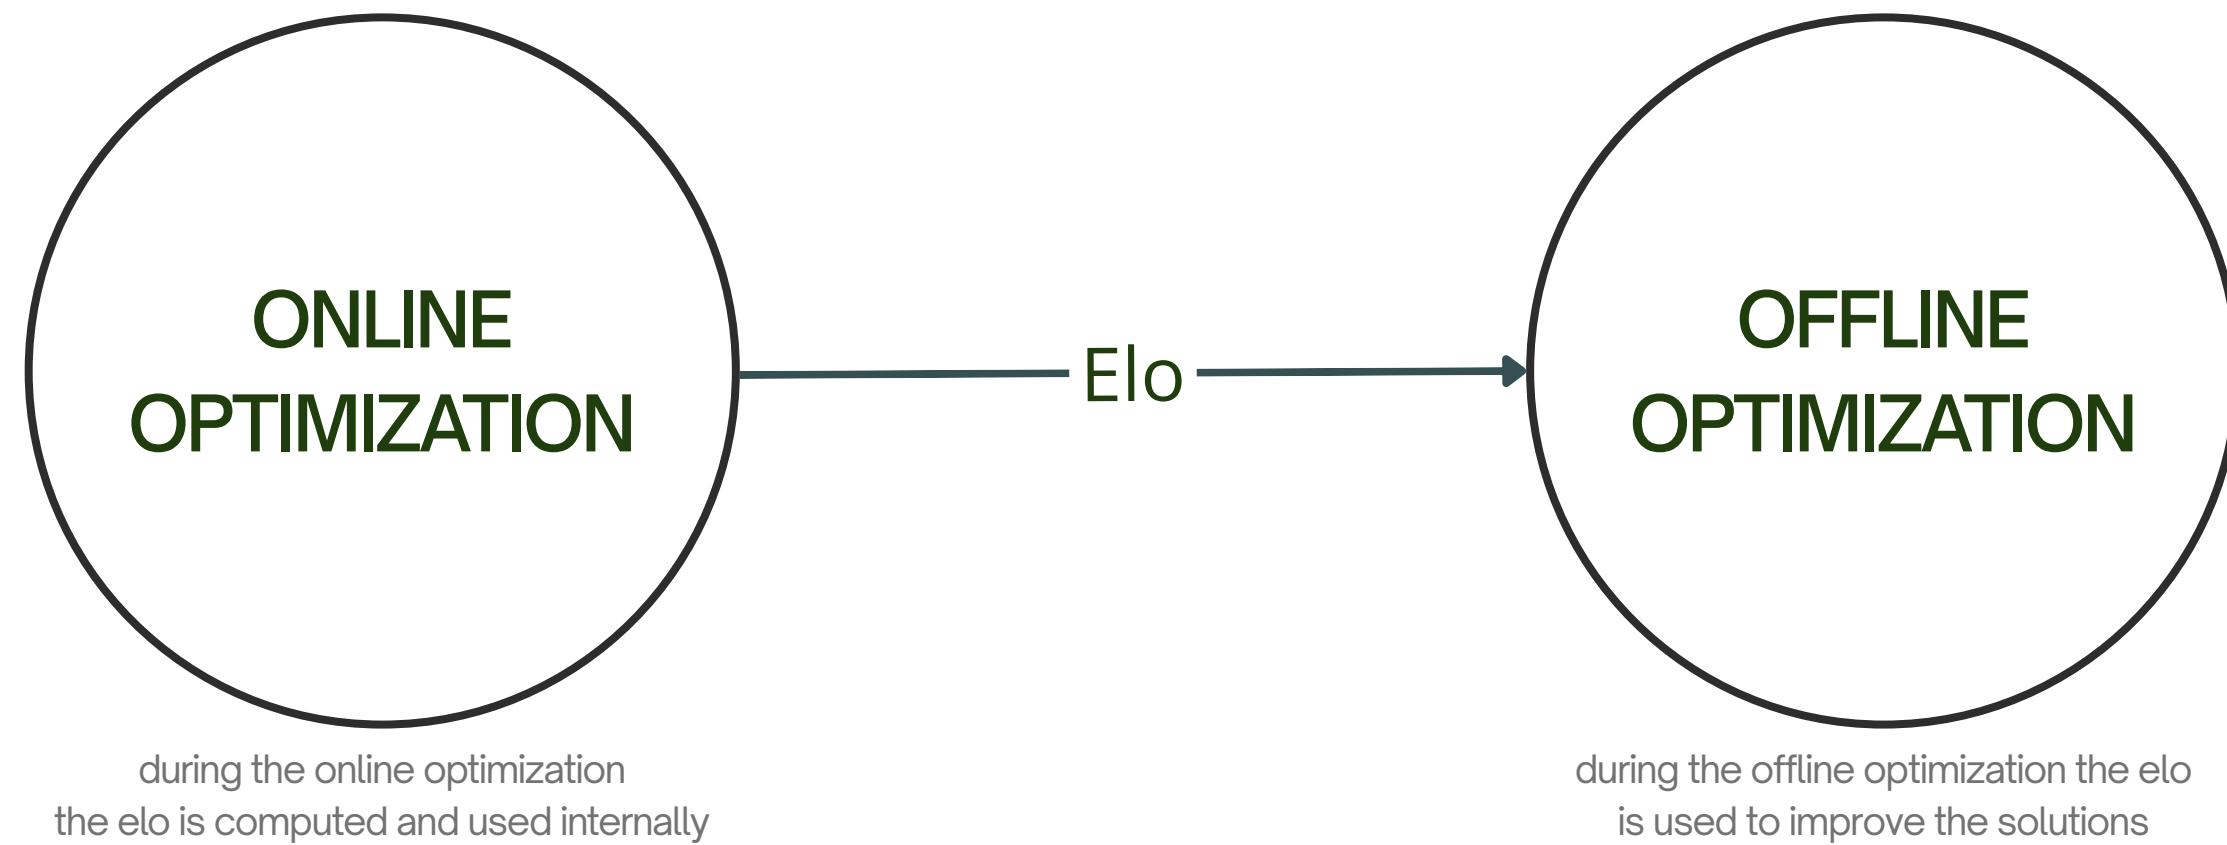
Reward signal approaches

reinforcement learning
learning classifier system

...

On the way to the EvES

Elo as link

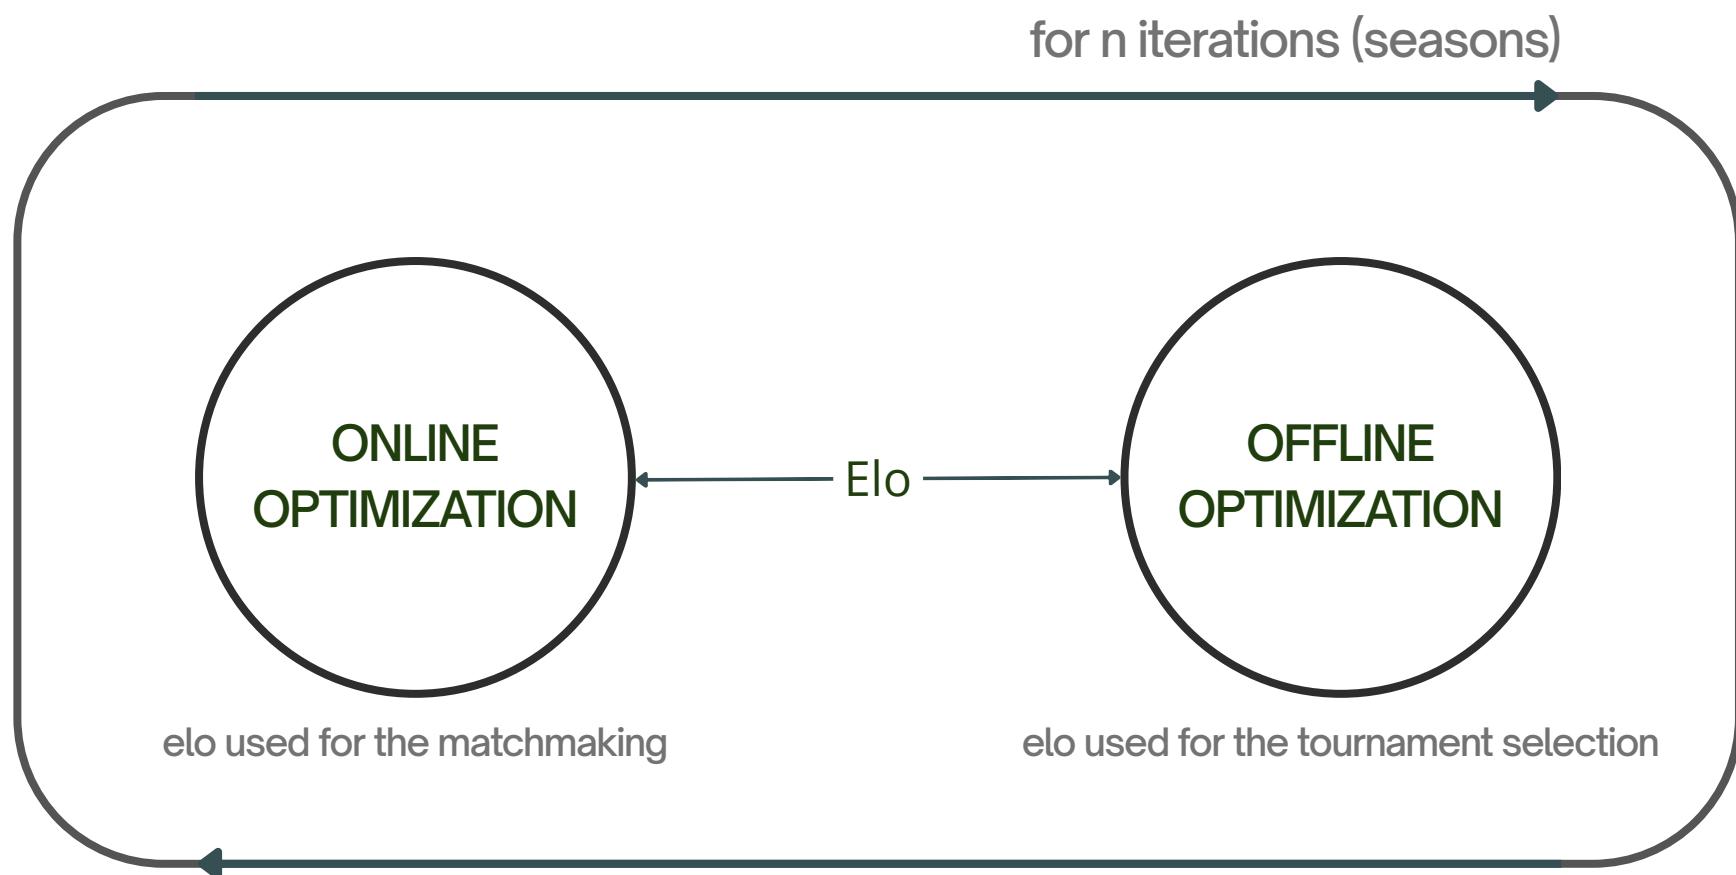


Bringing all together

EvES framework

It makes sense to develop an **hybrid method** to leverage the strengths of **both optimization paradigms**

In this framework individuals **must follow both protocols at the same time**



Key components

Environment

The environment must send signals to individuals in order to make the online optimization possible

A good approach is to follow the Gym/ALE protocol

Individuals are indirectly part of the environment and this is indeed fundamental

```
env = Env()
obs, info = env.reset()
done, truncate = False, False
while not done or truncated:
    actions = individuals.move(obs)
    new_obs, rewards, done, truncated, info = env.step(actions)
    players.observe(obs, actions, new_obs, rewards, done)
    obs = new_obs
env.close()
```

Key components

Individuals

Individuals must implement a variety of methods in order to accomplish both online and offline optimization

In order to produce good individuals one must ensure the capability of producing such individuals from the online optimization

Reactivity is a good feature

```
class Individual:  
    def observe(self, obs, action, rew, new_obs, done, **kwargs):  
        pass  
  
    def move(self, **kwargs):  
        pass  
  
    def reset(self, **kwargs):  
        pass  
  
    def update(self, **kwargs):  
        pass  
  
    def mutate(self, **kwargs):  
        pass
```

Minimal requirements for an individual

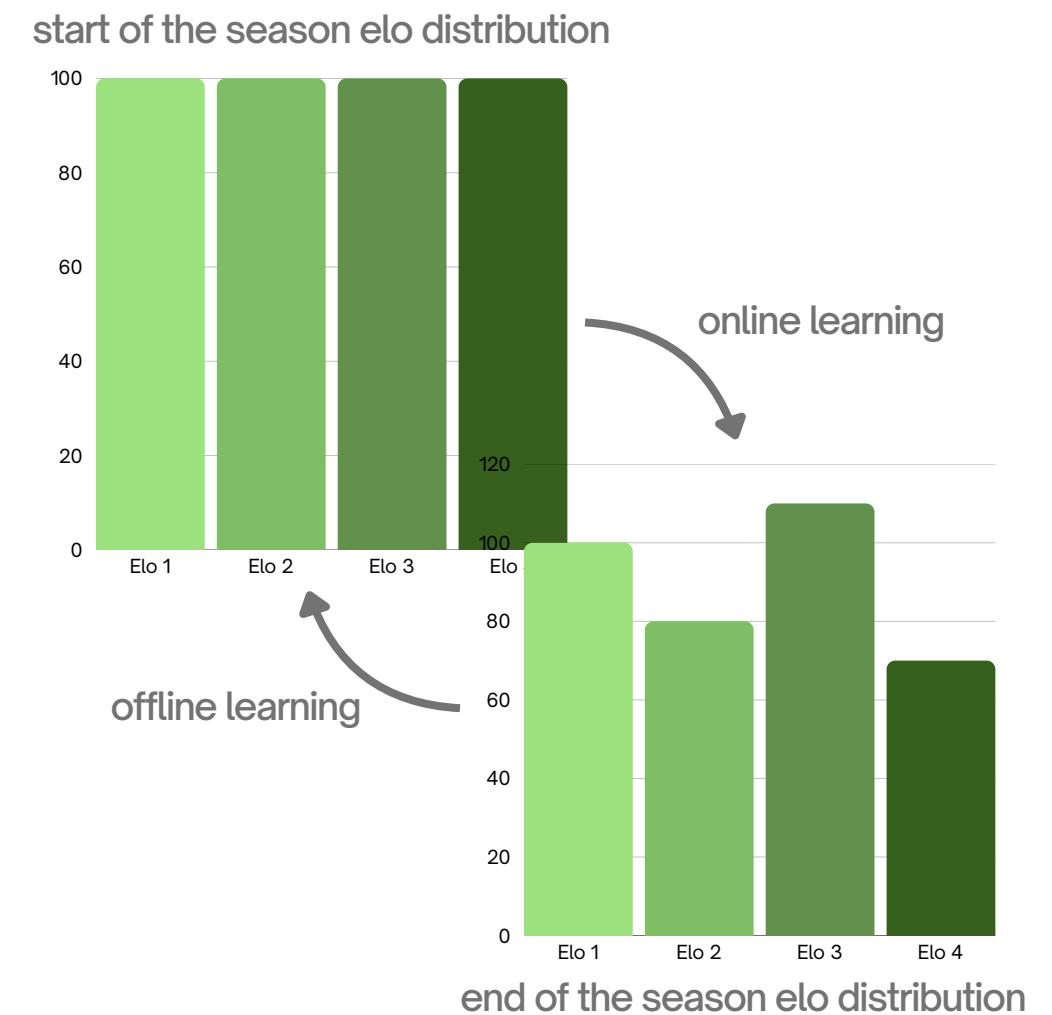
Key components

Elo Ranking System

An Elo Ranking System must be implemented inside the framework

It can just follow everything described up until now

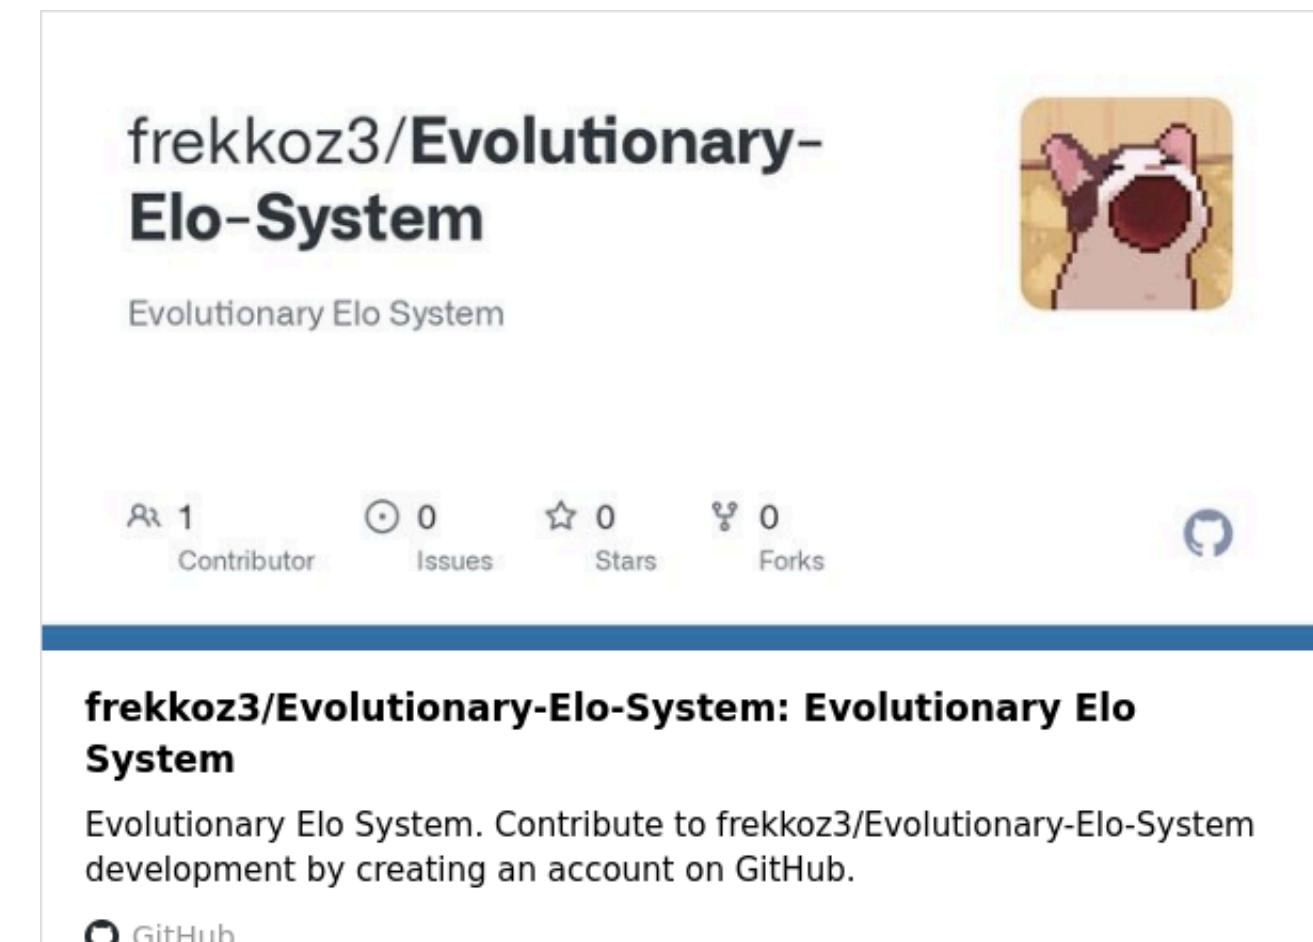
The elo of individuals should be reset at the beginning of a new season to a fixed level since there are no way to map the effect of a mutation to the Elo



About the Python implementation

The framework described until now is been implemented in **Python**

Several components are mappable 1:1 from the idea to the source code, while others have a variety of details better explained in their own README files

A GitHub repository card for the 'Evolutionary-Elo-System' project. The card includes the repository name, a small profile picture of a dog, and statistics: 1 contributor, 0 issues, 0 stars, and 0 forks. A blue header bar contains the repository name again. Below the bar, a description states: 'Evolutionary Elo System. Contribute to frekkoz3/Evolutionary-Elo-System development by creating an account on GitHub.' A GitHub logo is at the bottom.

frekkoz3/**Evolutionary-Elo-System**

Evolutionary Elo System

1 Contributor 0 Issues 0 Stars 0 Forks

frekkoz3/Evolutionary-Elo-System: Evolutionary Elo System

Evolutionary Elo System. Contribute to frekkoz3/Evolutionary-Elo-System development by creating an account on GitHub.

 GitHub

Details

DQN

The main online optimization techniques that I decided to use is the DQN, a technique that comes from the Reinforcement Learning area

It is a powerful technique but there are many caveats one should be aware of

A better choice could be the R2D2, capable of learning more complex dependencies



Details

Evolutionary Strategy

I decided to use a simple ES with elitism

There are no crossovers involved in the evolutionary cycle

The mutations are uniformly random

The tournament selection is the default selection type



Details

Benchmarks

Boxing

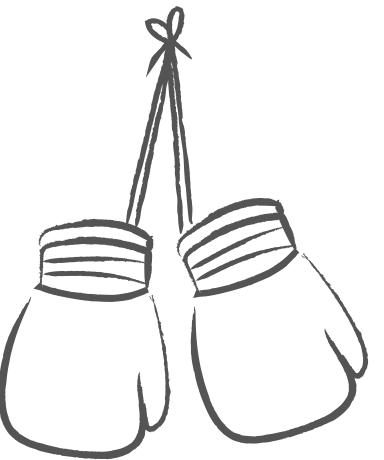
Rectangular ring

8 types of action

Stamina resource

Complex internal dynamic

The first to hit 15 punches wins



Grab 'n go

Rectangular ring

5 types of action

Maximum time 30 seconds

Problem in the Benchmarks

Boxing

The framework is too computational heavy

DQN does not capture time relationships

Individuals are way too few reagents

Elo to fine-tune

Grab 'n go

There are problem with the DQN

A really naive approach with a
Lamarckian GP didn't produce any results
(LCS would have been better)

The Elo is not really necessary

Conclusion

Take home lecture

This could have been a really interesting optimization technique but there are a lot of caveaut that make it really hard to work properly

Combining this technique with proper distributed methods could solve some of the main problems

Focusing on the online optimization is crucial in order to obtain good results

The Elo can be a bad measurement under certain conditions and other metrics should be used instead (Glicko, TrueSkill2, BayesElo)

Thank for your attention

Bredariol Francesco